

Multithreading

Lecture 2

Topic Outline


- Concepts of tasks, threads and multithreading.
- Thread pools and executors.
- Thread synchronization.
- Thread cooperation.

Multithreading


- Multithreading
 - powerful feature of Java
- Basic thread concept:
 - A **thread** is a flow of execution, from beginning to end, of a task in a program.
 - A **task** is a program unit that is executed independently of other parts of the program.
 - A **thread** provides the mechanism for running a **task**.

Concurrent Execution of Threads

- Multiple threads on multiple CPUs



- Multiple threads sharing a single CPU
 - Multiple threads share CPU time. The operating system is responsible for scheduling and allocating resources to threads.



Introduction to Multithreading

- Multithreading can make your program more responsive and interactive as well as enhance performance
- When a program executes as application, the Java interpreter starts a thread for the main method.
- In Java, each task is an instance of the **Runnable** interface, also called a runnable object. A thread is essentially an object that facilitates the execution of a task.

Creating Tasks & Threads

- Tasks are objects.
- To create tasks, you have to first declare a class of tasks. A task class must implement the **Runnable** interface.
 - **Runnable** interface contains a **run** method.
 - Once you have declared a task class class, you can create a task using its constructor

```
TaskClass task = new TaskClass();
```

Task Class

```
public class TaskClass implements Runnable {  
    ...  
        //constructor  
    public TaskClass() {  
        ...  
    }  
  
        //implement the run method in Runnable  
    public void run() {  
        // tell the system how to run  
        // custom thread  
        ...  
    }  
    ...  
}
```

Creating Tasks & Threads

- A task must be executed in a method. The **Thread** class contains the constructor for creating threads and many useful methods for controlling threads.

```
Thread thread = new Thread(task) ;
```

- You can then invoke the **start()** method to tell the JVM that the thread is ready to run, as follows:

```
thread.start() ;
```

- The JVM will execute the task by invoking the task's **run()** method

Thread Class

- The **Thread** class contains the constructors for creating threads for tasks, and the methods for controlling threads

«interface»
java.lang.Runnable



java.lang.Thread

+Thread()
+Thread(task: Runnable)
+start(): void
+isAlive(): boolean
+setPriority(p: int): void
+join(): void
+ <u>sleep(millis: long): void</u>
+ <u>yield(): void</u>
+interrupt(): void

Creates a default thread.

Creates a thread for a specified task.

Starts the thread that causes the run() method to be invoked by the JVM.

Tests whether the thread is currently running.

Sets priority p (ranging from 1 to 10) for this thread.

Waits for this thread to finish.

Puts the runnable object to sleep for a specified time in milliseconds.

Causes this thread to temporarily pause and allow other threads to execute.

Interrupts this thread.

Thread Class

- Since the **Thread** class implements **Runnable**, you could declare a class that extends `Thread` and implements `run()` method.
- Then, create an object from the class and invoke its `start` method in a client program to start the thread.

yield() Method

- You can use the **yield()** method to temporarily release time for other threads.
- In the **TaskThreadDemoYield.java** program:
 - Every time a number is printed, the **print100** thread is **yielded**.
 - So, the numbers are printed after the characters.

sleep() Method

- The **sleep(long)** method is used to put the thread to sleep for the specified time in milliseconds to allow other threads to execute.
- The **sleep(long)** method may throw an **InterruptedException**, which is a checked exception.
- This exception may occur when a sleeping thread's **interrupt()** method is called.
- We usually don't interrupt the thread so **InterruptedException** is unlikely to occur.

join() Method

- The **join()** method is used to force one thread to wait for another thread to finish
- In the **TaskThreadDemoJoin.java** program:
 - The numbers after 25 are printed after thread **printY** is finished.

Thread Priority

- By default, a thread inherits the priority of the thread that produced it.
- You can increase or decrease the priority of any thread by using the **setPriority()** method.
- You can also get the thread's priority by using **getPriority()** method.
- Priorities are numbers ranging from 1 to 10.
- There are constant like **MIN_PRIORITY**, **NORM_PRIORITY** and **MAX_PRIORITY**, representing 1, 5, and 10.

Thread Priority

- The JVM always picks the currently runnable thread with the highest priority.
- If several runnable threads have equally high priority, the CPU is allocated to all in round-robin fashion.
- Of course, a low priority thread can only be run when no higher-priority threads are running.


TaskThreadDemoPriority.java

Thread Pools

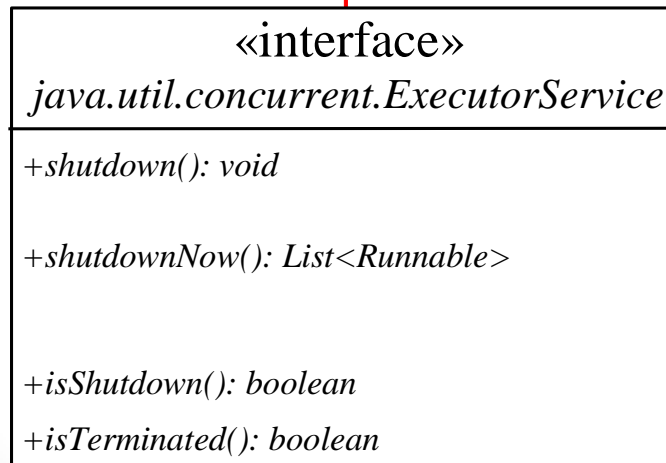
- You have seen how to declare a task class by implementing **Runnable**, and how to create a thread to run a task.
- This approach is convenient for a single task execution, but it isn't efficient for large number of tasks.

Thread Pools

- A **thread pool** is ideal to manage the number of tasks executing concurrently.
- Java uses the **Executor** interface for executing tasks in a thread pool and the **ExecutorService** interface for managing and controlling tasks.



Executes the runnable task.



Shuts down the executor, but allows the tasks in the executor to complete. Once shutdown, it cannot accept new tasks.

Shuts down the executor immediately even though there are unfinished threads in the pool. Returns a list of unfinished tasks.

Returns true if the executor has been shutdown.

Returns true if all tasks in the pool are terminated.

Creating Executors

- To create an **Executor** object, use the static methods in the **Executors** class.

java.util.concurrent.Executors
+newFixedThreadPool(numberOfThreads: int): ExecutorService
+newCachedThreadPool(): ExecutorService

Creates a thread pool with a fixed number of threads executing concurrently. A thread may be reused to execute another task after its current task is finished.

Creates a thread pool that creates new threads as needed, but will reuse previously constructed threads when they are available.

ExecutorDemo.java

Thread Synchronization

- A shared resource may be corrupted if it is accessed simultaneously by multiple threads.
- The program **AccountWithoutSync.java** demonstrate the data corruption problem as all the threads have access to the same data source simultaneously.
- You will see the result is unpredictable.

Thread Synchronization

- In **AccountWithoutSync.java**, Thread 1 and Thread 2, etc are accessing a common resource in a way that causes conflict.
- This is common problem, known as *race condition* in multithreaded program.
- A class is said to be thread-safe if an object of the class does not cause a race condition in the presence of multiple threads.
- Account class is not thread-safe in the example.

Thread Synchronization

- To avoid race condition, it is necessary to prevent more than one thread from simultaneously entering a certain part of the program (known as *critical region*)
- The critical region in **AccountWithoutSync.java** is **deposit()** method.
- Use the keyword **synchronized** to synchronize the method so that only one thread can access the method at a time.

synchronized Method

- A synchronized method acquires a lock before it executes.
- There are 2 cases:
 - The instance method, the lock is on the object when its invoked.
 - The static method, the lock is on the class.

synchronized Method

- If a thread invokes a synchronized instance method on an object, the lock of that object is acquired first, then the method is executed, and finally the lock is released.
- Another thread invoking the same method of that object is blocked until the lock is released.
- Put the synchronized keyword on the **deposit** method as follows:

```
public synchronized void deposit(int amount) {...}
```

AccountWithSyncMethod.java

synchronized Statement

- A synchronized statement can be used to acquire a lock on any object, not just this object, when executing a block of the code in a method.
- This is referred to as a synchronized block.

```
synchronized (expr) {  
    statements;  
}
```


synchronized Statement

- The expression **expr** must evaluate to an object reference.
- If the object is already locked by another thread, the thread is blocked until the lock is released.
- When a lock is obtained on the object, the statements in the synchronized block are executed, then the lock is released.

synchronized Statement

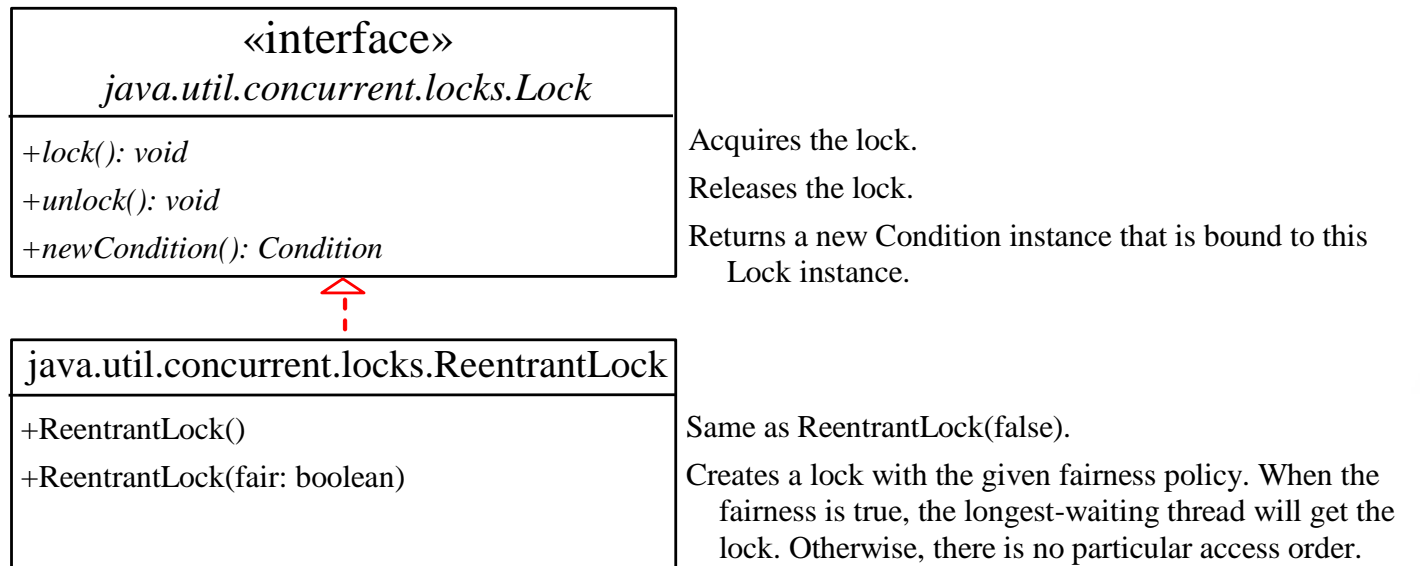
- Synchronized statements enable you to synchronize part of the code in a method instead of the entire method.
 - This increases concurrency.
- Synchronized statements enable you acquire a lock on any object so that you can synchronize the access to an object instead of to a method.

```
synchronized (account) {  
    account.deposit(1) ;  
}
```

AccountWithSyncStatement.java

Synchronization Using Locks

- A synchronized instance method implicitly acquires a lock on the instance before it executes the method.
- With the use of locks explicitly, the locking features are flexible and give you more control for coordinating threads.



Thread Cooperation

- Sometimes, you need a way for threads to cooperate.
- Conditions can be used to facilitate communications among threads. The thread specify what to do under a certain condition.
- Once condition is created, you can use its **`await()`**, **`signal()`**, and **`signalAll()`** methods for thread communications.

Thread Cooperation

«interface»

java.util.concurrent.Condition

+*await(): void*

+*signal(): void*

+*signalAll(): Condition*

Causes the current thread to wait until the condition is signaled.

Wakes up one waiting thread.

Wakes up all waiting threads.

Thread Cooperation

- To synchronize the operations, use a lock with a condition.
- Example:
 - Condition: **newDeposit** (i.e., new deposit added to the account).
 - If the balance is less than the amount to be withdrawn, the withdraw task will wait for the **newDeposit** condition.
 - When the **deposit** task adds money to the account, the task signals the waiting withdraw task to try again.

Thread Cooperation

